

Hardware Overview

Andreas Götz

Material by
Robert Sinkovits,
Mary Thomas,
Mahidhar Tatineni

SDSC Summer Institute
August 1, 2022



Introduction

- Most of you are here because you are computational scientists
 - Have a specific scientific problem you're trying to solve
 - Support researchers from a variety of domains
- The actual hardware is probably of secondary interest
 - Hardware is interesting but not *that* interesting
- Nonetheless, it's still helpful to know a bit about hardware
- In this talk we will learn
 - Expanse hardware
 - Getting information about your system
 - Using some common usage monitoring tools

Outline

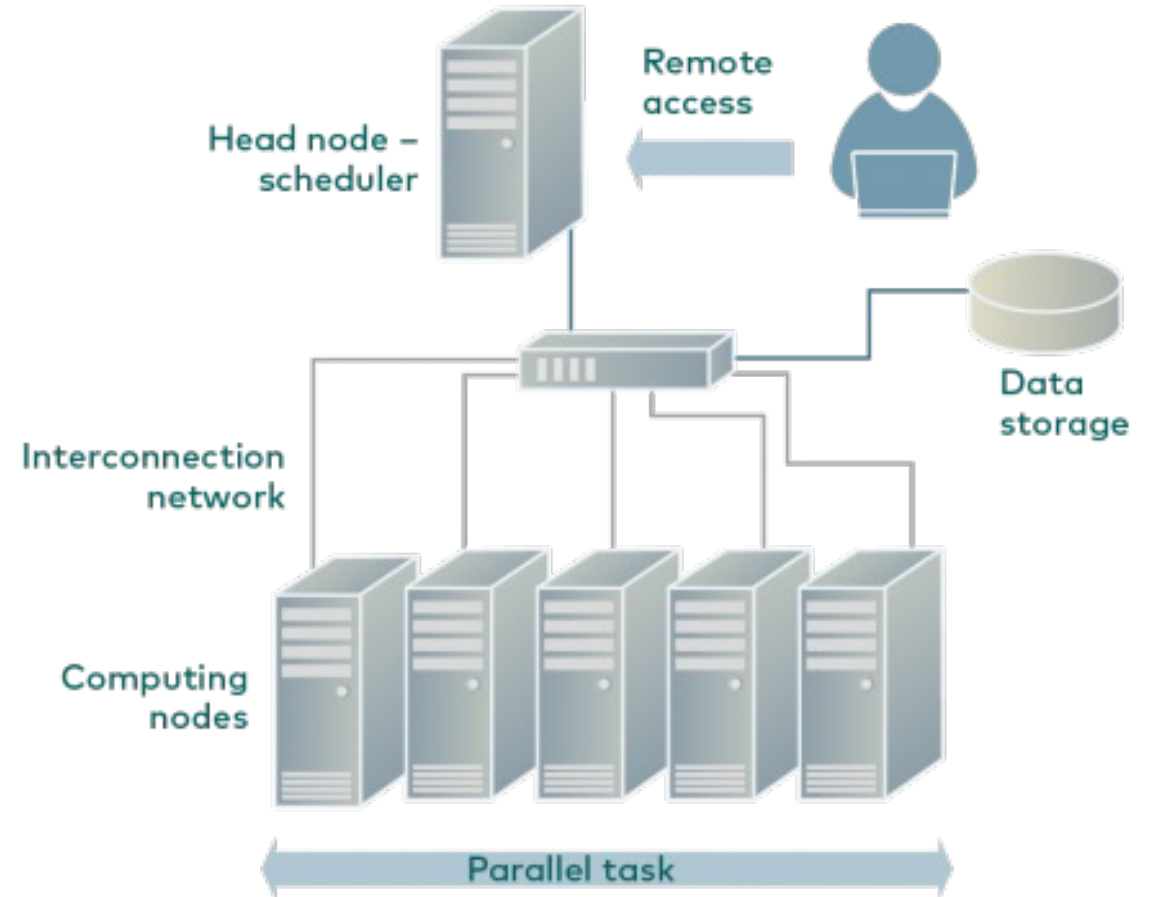
- Expanse
 - System Overview
 - AMD EPYC Processor Architecture
- Obtaining Hardware Information
 - CPU
 - GPU
 - Memory
 - Cache
 - SCSI, df, network, OS
- top and htop tools

Outline

- Expanse
 - **System Overview**
 - AMD EPYC Processor Architecture
- Obtaining Hardware Information
 - CPU
 - GPU
 - Memory
 - Cache
 - SCSI, df, network, OS
- top and htop tools

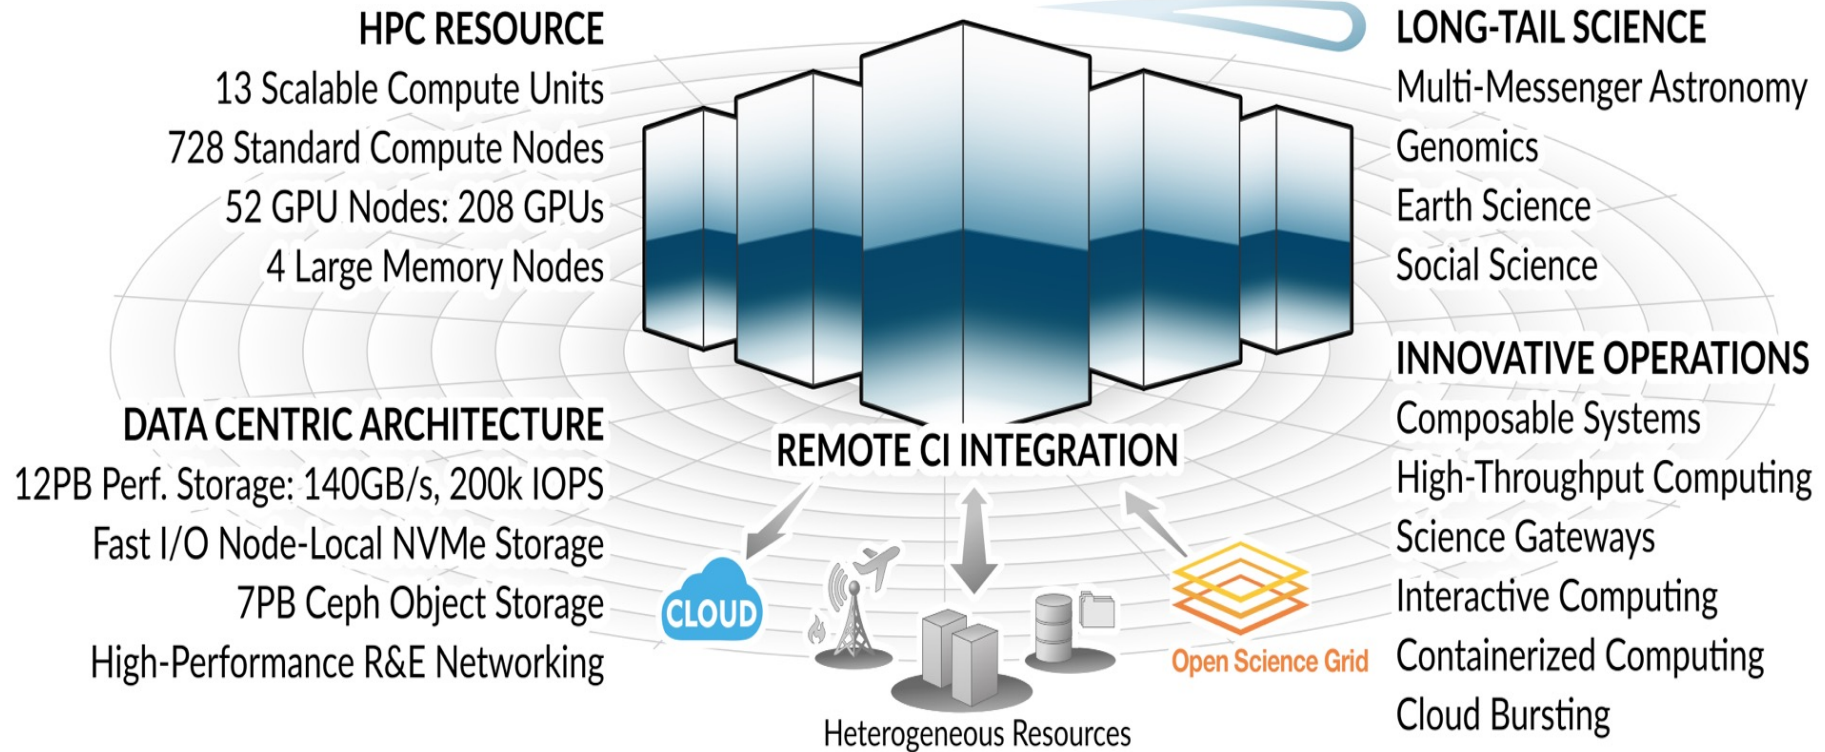
System overview

- Linux/Mac –
 - use terminal + installed ssh app
- Windows:
 - Win10 terminal app + installed ssh app
 - Older Windows OS's: ssh clients apps Putty, Cygwin
- Login hostname for SDSC Expanse:
 - login.expense.sdsc.edu
 - 198.202.113.252



EXPANSE

COMPUTING WITHOUT BOUNDARIES
5 PETAFLOP/S HPC and DATA RESOURCE



For more details see the Expanse user guide @ https://www.sdsc.edu/support/user_guides/expanse.html
and the "Introduction to Expanse" webinar @ https://www.sdsc.edu/event_items/202006_Introduction_to_Expanse.html

Expanse



SDSC

SAN DIEGO SUPERCOMPUTER CENTER

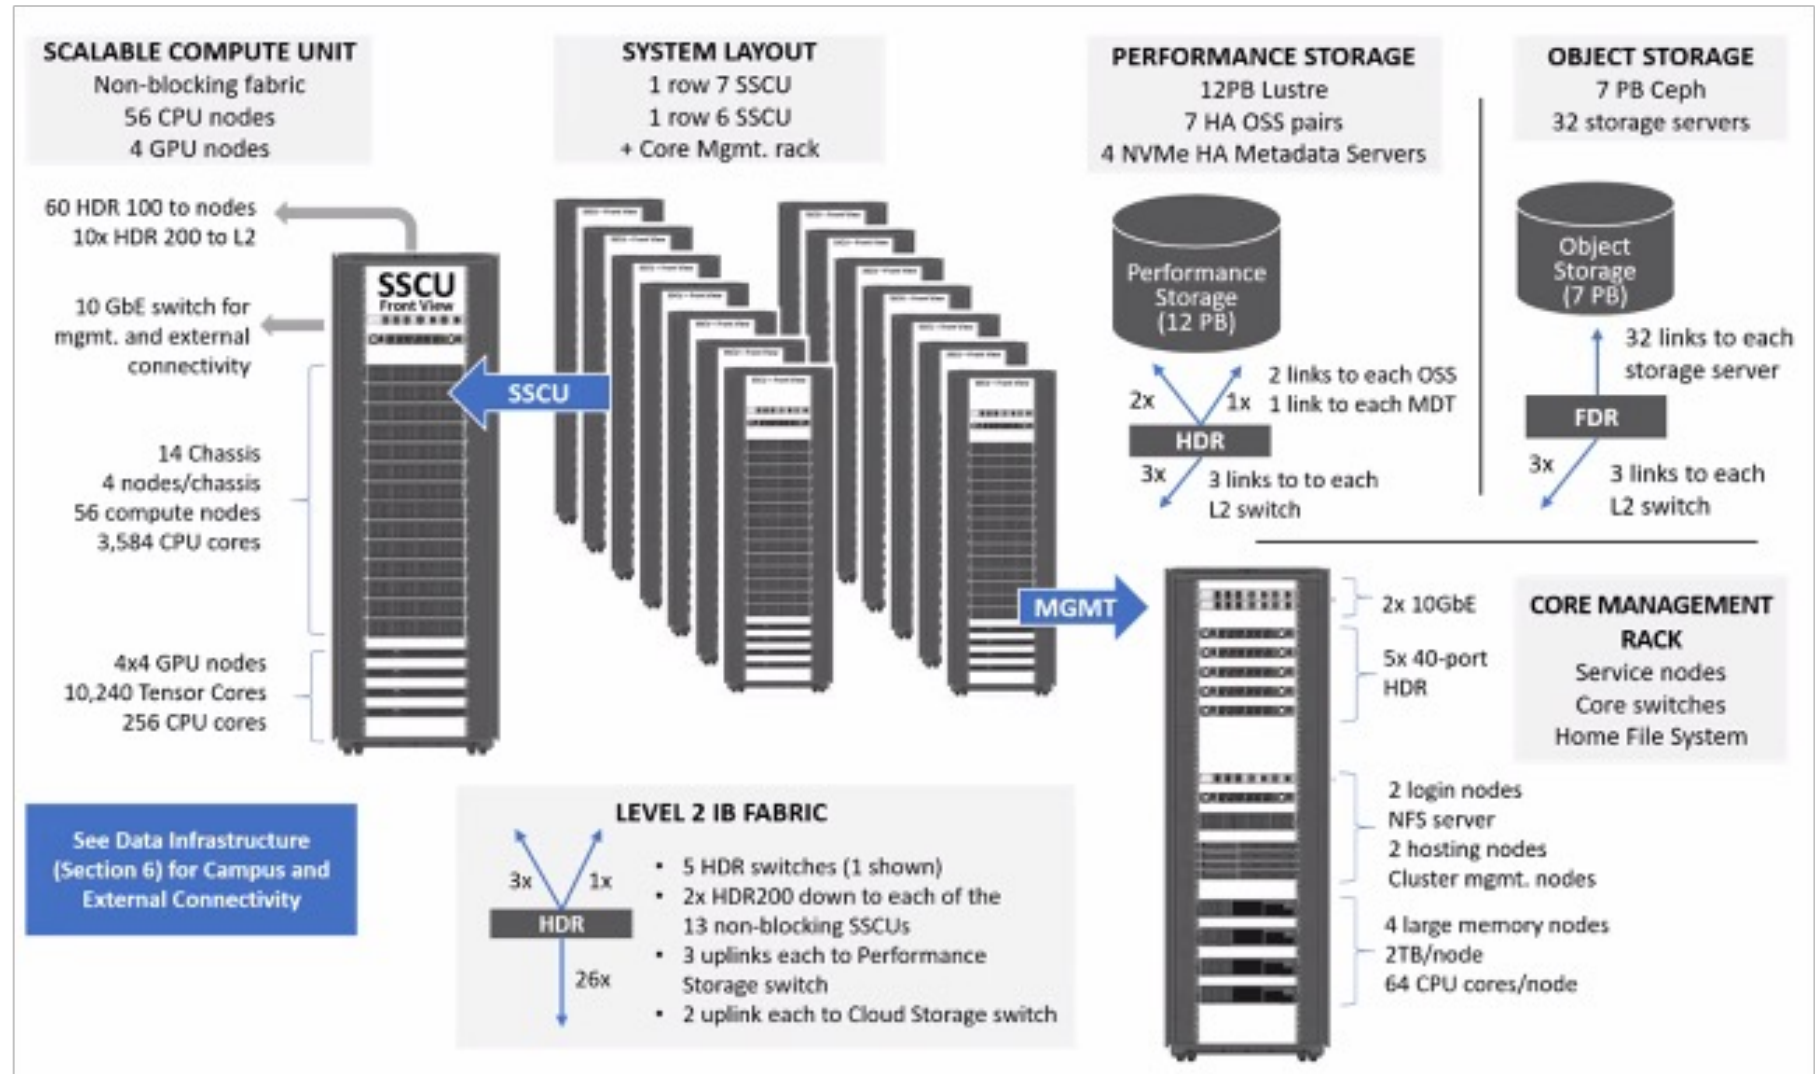
at the UNIVERSITY OF CALIFORNIA; SAN DIEGO



Expanse Heterogeneous Architecture

System Summary

- 13 SDSC Scalable Compute Units (SSCU)
- 728 x 2s Standard Compute Nodes
- 93,184 Compute Cores
- 200 TB DDR4 Memory
- 52x 4-way GPU Nodes w/NVLINK
- 208 V100s
- 4x 2TB Large Memory Nodes
- HDR 100 non-blocking Fabric
- 12 PB Lustre High Performance Storage
- Storage
- 7 PB Ceph Object Storage
- 1.2 PB on-node NVMe
- Dell EMC PowerEdge
- Direct Liquid Cooled



Using Login Nodes Properly

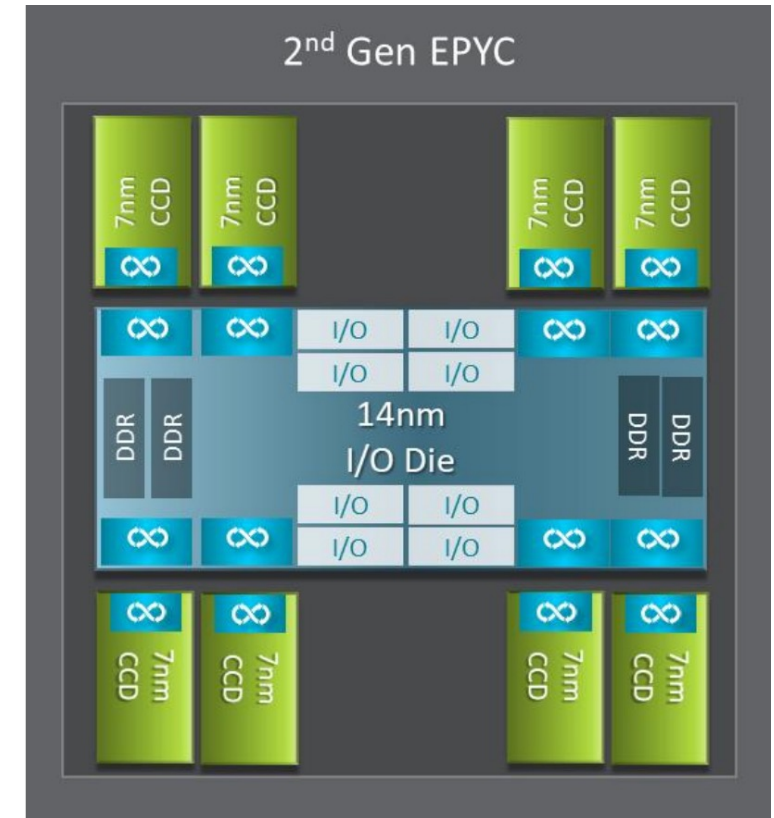
- The login nodes are meant for file editing, simple data analysis, & tasks that use minimal compute resources.
- All computationally demanding jobs should be submitted and run through the batch queuing system.
- **Do not use the login nodes for:**
 - computationally intensive processes,
 - as hosts for running workflow management tools
 - as primary data transfer nodes for large or numerous data transfers
 - as servers providing other services accessible to the Internet.
 - running Jupyter notebooks
- **Login nodes are not the same as the batch nodes.**
 - Users should request an interactive sessions to compile large programs.

Outline

- Expanse
 - System Overview
 - **AMD EPYC Processor Architecture**
- Obtaining Hardware Information
 - CPU
 - GPU
 - Memory
 - Cache
 - SCSI, df, network, OS
- top and htop tools

AMD EPYC 7742 Processor Architecture

- 8 Core Complex Dies (CCDs).
- CCDs connect to memory, I/O, and each other through the I/O Die.
- 8 memory channels per socket.
- DDR4 memory at 3200MHz.
- PCI Gen4, up to 128 lanes of high speed I/O.
- Memory and I/O can be abstracted into separate quadrants each with 2 DIMM channels and 32 I/O lanes.



Reference: <https://developer.amd.com/wp-content/resources/56827-1-0.pdf>

AMD EPYC 7742 Processor Architecture

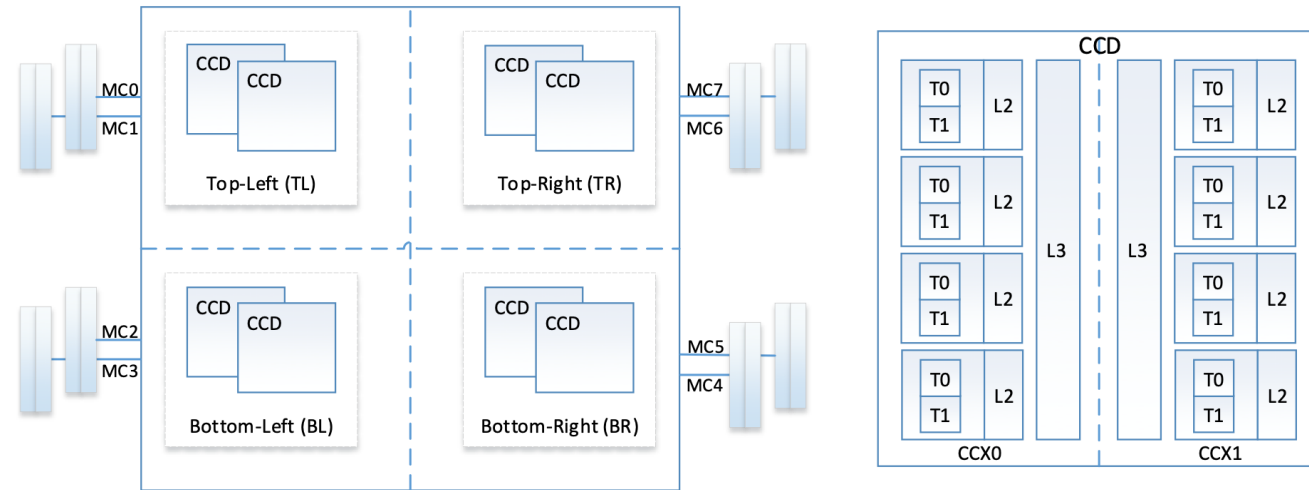
- 2 Core Complexes (CCXs) per CCD
- 4 Zen2 cores in each CCX shared a 16M L3 cache. Total of $16 \times 16 = 256\text{MB}$ L3 cache.
- Each core includes a private 512KB L2 cache.



Reference: <https://developer.amd.com/wp-content/resources/56827-1-0.pdf>

AMD EPYC 7742 Processor Architecture


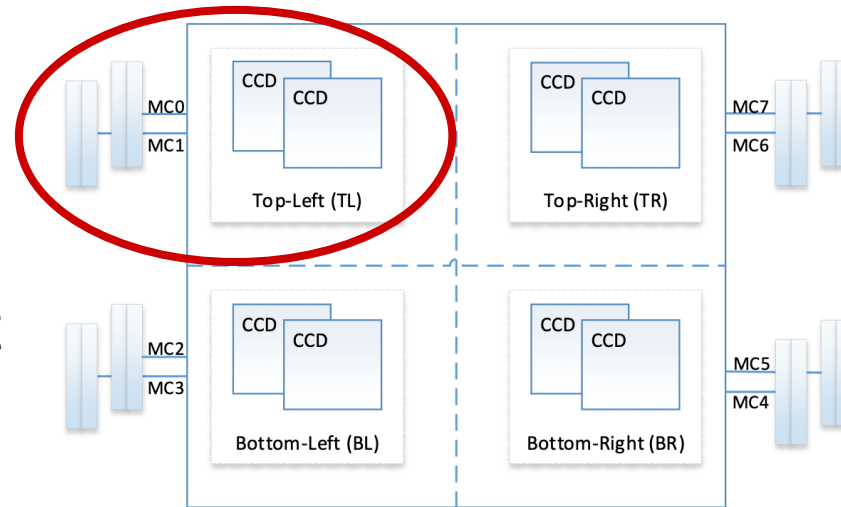
- The four logical quadrants allow the processor to be partitioned into different NUMA domains. Options set in BIOS.
- Domains are designated as NUMA per socket (NPS).
- **NPS4:** Four NUMA domains per socket is the typical HPC configuration.



https://developer.amd.com/wp-content/resources/56338_1.00_pub.pdf

AMD EPYC 7742 Processor Architecture

- The processor is partitioned into four NUMA domains.
- Each logical quadrant is a NUMA domain.
- Memory is interleaved across the two memory channels
- PCIe devices will be local to one of four NUMA domains (the IO die that has the PCIe root for the device)
- ***This is the typical HPC configuration*** as workload is NUMA aware, ranks and memory can be pinned to cores and NUMA nodes.



https://developer.amd.com/wp-content/resources/56338_1.00_pub.pdf

Expanse Compute Node

The AMD processors on Expanse

- Have 64 cores in total.
- Consist of 8 core complex dies (CCDs) per processor, with 2 core complexes (CCXs) per CCD. Four cores in a CCX share L3 cache.
- There are 4 NUMA domains per processor.

Each Expanse compute node

- Has 2 AMD processors, for a total of 128 cores.
- Has 2 GB RAM per core, for a total of 256 GB RAM.
- Has 1 TB NVMe storage (for fast scratch I/O)

Compute nodes are interconnected with HDR Infiniband.

Expanse storage

It is important to understand the storage hierarchy on Expanse (more later).

- The home file system is mounted via NFS and used to store smaller amounts of important data like source code files. This is slow I/O storage, not suitable for heavy I/O when running jobs.
- The Lustre file system is suitable for parallel I/O and large amount of project files like simulation output data.
- The local NVMe (SSD) storage is limited in space but very fast and should be used when doing heavy I/O during simulations.

Outline


- Expanse
 - System Overview
 - AMD EPYC Processor Architecture
- Obtaining Hardware Information
 - CPU
 - GPU
 - Memory
 - Cache
 - SCSI, df, network, OS
- top and htop tools

Getting hardware information – why do I care?

- You may be asked to report details of your hardware in a manuscript, presentation, proposal or request for computer time
- You'll know what you're running on and can answer questions like
 - Is the login node the same as the compute nodes?
 - How does one machine compare to another?
- It will give you a way of estimating performance or at least bounds on performance relative to another system. *All else being equal*, jobs will run at least as fast on hardware with
 - Faster CPU clock speeds
 - Larger caches
 - Faster local drives

Computers are like hammers!


Just like hammers, there is a wide variety of computer hardware. You can probably get away with using the wrong one, but your performance may be suboptimal and you might end up using a bigger tool than you need.



Ball-peen hammer



Ball peen hammer



Bush hammer




Claw hammer




Cross-peen hammer




Dog-head hammer
(blacksmithing)



Framing hammer




Geologist's hammer




HiFIT-hammer for
aftertreatment of weld
transitions



Long cross-face hammer
(blacksmithing)




Post maul




Rock climbing hammer



Rubber mallet




Sledgehammers



Straight pane
sledgehammer



Twist hammer
(blacksmithing)



Upholstery hammer



Wooden mallet

Wikipedia

Processor specifications: lscpu

On Linux systems, the **lscpu** command lists key processor information

- Number of processors (sockets)
- Processor type or model
- Nominal clock speed
- Number of cores per processor
- Cache sizes
- Instruction set architecture
- NUMA nodes

Processor specifications: Expanse compute node

Architecture: x86_64
CPU op-mode(s): 32-bit, 64-bit
Byte Order: Little Endian
CPU(s): 128
On-line CPU(s) list: 0-127
Thread(s) per core: 1
Core(s) per socket: 64
Socket(s): 2
NUMA node(s): 8
Vendor ID: AuthenticAMD
CPU family: 23
Model: 49
Model name: AMD EPYC 7742 64-Core Processor
Stepping: 0
CPU MHz: 3257.493
BogoMIPS: 4491.71
Virtualization: AMD-V
L1d cache: 32K
L1i cache: 32K
L2 cache: 512K
L3 cache: 16384K

NUMA node0 CPU(s): 0-15
NUMA node1 CPU(s): 16-31
NUMA node2 CPU(s): 32-47
NUMA node3 CPU(s): 48-63
NUMA node4 CPU(s): 64-79
NUMA node5 CPU(s): 80-95
NUMA node6 CPU(s): 96-111
NUMA node7 CPU(s): 112-127
Flags: fpu vme de pse tsc msr pae mce
cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx
fxsr sse sse2 ht syscall nx mmxext fxsr_opt pdpe1gb
rdtscp lm constant_tsc rep_good nopl xtopology
nonstop_tsc cpuid extd_apicid aperfmperf pni
pclmulqdq monitor ssse3 fma cx16 sse4_1 sse4_2
x2apic movbe popcnt aes xsave avx f16c rdrand
lahf_lm cmp_legacy svm extapic cr8_legacy abm sse4a
misalignsse 3dnowprefetch osvw ibs skinit wdt tce
topoext perfctr_core perfctr_nb bpext perfctr_llc
mwaitx cpb cat_l3 cdp_l3 hw_pstate sme ssbd mba sev
ibrs ibpb stibp vmmcall fsgsbase bmi1 avx2 smep bmi2
cqm rdt_a rdseed adx smap clflushopt clwb sha_ni
xsaveopt xsavec xgetbv1 xsaves cqm_llc cqm_occup_llc
cqm_mbm_total cqm_mbm_local clzero irperf xsaveerptr
wbnoinvd arat npt lbrv svm_lock nrip_save tsc_scale
vmcb_clean flushbyasid decodeas

Processor specifications: Expanse compute node

Architecture: x86_64
CPU op-mode(s): 32-bit, 64-bit
Byte Order: Little Endian
CPU(s): 128
On-line CPU(s) list: 0-127
Thread(s) per core: 1
Core(s) per socket: 64
Socket(s): 2
NUMA node(s): 8
Vendor ID: AuthenticAMD
CPU family: 23
Model: 49
Model name: AMD EPYC 7742 64-Core Processor
Stepping: 0
CPU MHz: 3257.493
BogoMIPS: 4491.71
Virtualization: AMD-V
L1d cache: 32K
L1i cache: 32K
L2 cache: 512K
L3 cache: 16384K

NUMA node0 CPU(s): 0-15
NUMA node1 CPU(s): 16-31
NUMA node2 CPU(s): 32-47
NUMA node3 CPU(s): 48-63
NUMA node4 CPU(s): 64-79
NUMA node5 CPU(s): 80-95
NUMA node6 CPU(s): 96-111
NUMA node7 CPU(s): 112-127
Flags: fpu vme de pse tsc msr pae mce
cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx
fxsr sse sse2 ht syscall nx mmxext fxsr_opt pdpe1gb
rdtscp lm constant_tsc rep_good nopl xtopology
nonstop_tsc cpuid extd_apicid aperfmperf pni
pclmulqdq monitor ssse3 fma cx16 sse4_1 sse4_2
x2apic movbe popcnt aes xsave avx fl6c rdrand
lahf_lm cmp_legacy svm extapic cr8_legacy abm sse4a
misalignsse 3dnowprefetch osvw ibs skinit wdt tce
topoext perfctr_core perfctr_nb bpext perfctr_llc
mwaitx cpb cat_l3 cdp_l3 hw_pstate sme ssbd mba sev
ibrs ibpb stibp vmmcall fsgsbase bmi1 avx2 smep bmi2
cqm rdt_a rdseed adx smap clflushopt clwb sha_ni
xsaveopt xsavec xgetbv1 xsaves cqm_llc cqm_occup_llc
cqm_mbm_total cqm_mbm_local clzero irperf xsaveerptr
wbnoinvd arat npt lbrv svm_lock nrip_save tsc_scale
vmcb_clean flushbyasid decodeas

Processor specifications: /proc/cpuinfo

On Linux systems, the **/proc/cpuinfo** pseudo-file contains pretty much the same information that you get from `lscpu`, but with a few differences

- Information is listed per core
- Access to instantaneous clock speeds
- Bugs detected / addressed
see discussion: <https://unix.stackexchange.com/questions/456425/what-does-the-bugs-section-of-proc-cpuinfo-actually-show>
- Microcode, TLB size , power management, cache line flush sizes and other low-level details that you probably don't need to know about

Processor specifications: /proc/cpuinfo


```
$ grep 'cpu MHz' /proc/cpuinfo | head -5
```

```
cpu MHz      : 3325.325  
cpu MHz      : 2239.263  
cpu MHz      : 3374.887  
cpu MHz      : 2360.551  
cpu MHz      : 2393.493
```

```
--- selected output ---
```

```
microcode     : 0x8301038  
bugs          : sysret_ss_attrs spectre_v1 spectre_v2 spec_store_bypass  
TLB size      : 3072 4K pages  
clflush size   : 64  
cache_alignment : 64  
address sizes  : 43 bits physical, 48 bits virtual  
power management : ts ttp tm hwpstate cpb eff_freq_ro [13] [14]
```

A brief aside on nomenclature



We normally think of the multicore unit that plugs into the motherboard as the “processor”

/proc/cpuinfo uses processor in a different way to mean compute core, counted across all the cores available in a node. lscpu uses more intuitive terminology.

processor	: 0
physical id	: 0
cpu cores	: 64

Quick aside on simultaneous multithreading

Simultaneous multithreading, abbreviated as SMT, is the process of a CPU splitting each of its physical cores into virtual cores, which are known as threads. This is done in order to increase performance and allow each core to run two instruction streams at once.

Intel branded this process as hyper-threading, but hyper-threading is the same thing as simultaneous multithreading. For example, AMD CPUs with four cores use simultaneous multithreading to provide eight threads, and most Intel CPUs with two cores use hyper-threading to provide four threads.

<https://www.tomshardware.com/reviews/simultaneous-multithreading-definition,5762.html>

SDSC does not enable hyperthreading on its systems. When hyperthreading is enabled, core count will appear to be doubled.

A brief aside on pseudo-files

Up to this point, we've been using the term pseudo-file without defining what it is. Recall that in the UNIX/Linux world, everything is treated as a file (files, directories, devices, etc.)

/proc and /sys are just interfaces to the Linux kernel data structures in a convenient and familiar file system format

```
$ ls -ld /proc
dr-xr-xr-x 2258 root root 0 Jul 28 09:27 /proc
[sinkovit@login01 ~]$ ls -ld /proc/cpuinfo
-r--r--r-- 1 root root 0 Jul 28 16:56 /proc/cpuinfo

$ head /proc/cpuinfo
processor      : 0
vendor_id     : AuthenticAMD
cpu family    : 23
Model.        : 49
model name    : AMD EPYC 7742 64-Core Processor
stepping      : 0
microcode     : 0x8301038
```

What's in a name?

We usually think of processors in terms of their codenames, such as Rome or Milan. Unfortunately, the `/proc/cpuinfo` pseudo-file returns something a little more opaque such as “AMD EPYC 7742 64-Core Processor”. A quick Google search helps

In November 2018 AMD announced Epyc 2 at their Next Horizon event, the second generation of Epyc processors code-named "**Rome**" and based on the **Zen 2** microarchitecture.^[19] The processors feature up to eight **7 nm**-based "chiplet" processors with a 14 nm-based **IO** chip providing 128 PCIe lanes in the center interconnected via **Infinity Fabric**. The processors support up to 8 channels of DDR4 RAM up to 4 **TB**, and introduce support for **PCIe** 4.0. These processors have up to 64 cores with 128 **SMT** threads per socket.^[20] The 7 nm "**Rome**" is manufactured by **TSMC**.^[11] It was released on August 7, 2019.^[21]


<https://en.wikipedia.org/wiki/Epyc>

Advanced Vector Extensions (AVX, AVX2, AVX512)

The Advanced Vector Extensions (AVX) are an extension to the x86 microprocessor architecture that allows a compute core to perform up to 8 floating point operations per cycle. Previous limit was 4/core/cycle

- AVX2 improves this to 16 Flops/cycle/core (Comet, Expanse)
- AVX512 further improves to 32 Flops/cycle/core (Intel \geq Skylake)

These were developed partially in response to challenges in increasing CPU clock speeds



<http://www.computerhope.com>

March 6, 2000 8:00 AM PST

AMD makes move to 1-GHz chip

By Joe Wilcox and Michael Kanellos
Staff Writers, CNET News

Advanced Vector Extensions (AVX, AVX2, AVX512)

- Can theoretically obtain a 2x speedup when going from a non-AVX processor to an AVX capable processor (all else being equal)
 - And another 2x from AVX to AVX2
 - And another 2x from AVX2 to AVX512
- But don't get too excited (or worried that Expanse doesn't have AVX512)
 - It's difficult enough to make good use of AVX and even harder to make good use of AVX2 or AVX512.
 - Need long loops with vectorizable content. Memory bandwidth not keeping up with gains in computing power.
 - On Skylake, clock speed scaled down when executing AVX512 instructions

Getting memory information: /proc/meminfo

On Linux machines, the /proc/meminfo pseudo-file lists key memory specs. More information than you probably want, but at least one bit of useful data

```
MemTotal:      263698228 kB (total physical memory)
MemFree:       251035032 kB
MemAvailable:  250623760 kB
Buffers:       12824 kB
Cached:        3126364 kB
SwapCached:    0 kB
Active:        1301564 kB (pretty good approximation to used memory)
Inactive:     2990668 kB
Active(anon):  1240284 kB
Inactive(anon): 2890076 kB
Active(file):  61280 kB
Inactive(file): 100592 kB
Unevictable:   0 kB
Mlocked:      0 kB
SwapTotal:     0 kB
SwapFree:      0 kB
Dirty:         32 kB
Writeback:     0 kB
AnonPages:     1151660 kB
```

Getting memory information (/proc/meminfo)

Using a simple script, you can monitor total memory usage for all processes as a function of time. Note that there is a lot of discussion on how to precisely measure memory (<http://stackoverflow.com/search?q=measuring+memory+usage>). The following should be good enough if you're on a dedicated node.

```
#!/usr/bin/perl
use strict;
use warnings;
my $count = 0;
print (" time(s)      Memory (GB)\n");
while(1) {
    sleep(1);
    $count++;
    open(MI, "/proc/meminfo");
    while(<MI>) {
        if (/Active:/) {
            my (undef, $active, undef) = split();
            $active = $active / 1048576.0;
            printf("%6d      %f\n", $count, $active);
        }
    }
    close(MI);
}
```

More memory information - dmidecode

If you really need to dig deeper and get more details on memory configuration, you can run the dmidecode command. You'll need root privileges to do this.

Output shows the results for one DIMM slot.

dmidecode --type memory

Memory Device

```
Array Handle: 0x001D
Error Information Handle: No Error
Total Width: 72 bits
Data Width: 64 bits
Size: 16384 MB
Form Factor: DIMM
Set: None
Locator: DIMM_A1
Bank Locator: CPU1
Type: DDR4
Type Detail: Synchronous Registered (Buffered)
Speed: 2133 MHz
Manufacturer: 0xCE00
Serial Number: 0x394FECDD
Asset Tag: Unknown
Part Number: M393A2G40DB0-CPB
Rank: 1
Configured Clock Speed: 2133 MHz
Minimum Voltage: 1.2 V
Maximum Voltage: 1.2 V
Configured Voltage: 1.2 V
```

Getting GPU information

If you're using GPU nodes, you can use nvidia-smi (NVIDIA System Management Interface program) to get GPU information (type, count, etc.)

When running in gpu-shared partition, will only see the GPUs you had requested (typically one)

```
$ nvidia-smi
Mon Aug  2 10:16:29 2021

+-----+
| NVIDIA-SMI 460.32.03      Driver Version: 460.32.03      CUDA Version: 11.2      |
+-----+-----+-----+
| GPU  Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                       |                    |     MIG M.     |
+-----+-----+-----+
| 0  Tesla V100-SXM2...  On      | 00000000:18:00.0 Off |                    0  |
| N/A   40C    P0     67W / 300W |      0MiB / 32510MiB |      0%      Default  |
|                                       |                    |     N/A     |
+-----+-----+-----+

+-----+
| Processes: |
| GPU   GI    CI          PID    Type    Process name                        GPU Memory |
|          ID    ID                                   |          Usage  |
+-----+-----+-----+
| No running processes found |
+-----+
```

Finding cache information

On Linux systems, can obtain cache properties through the /sys pseudo filesystem. Details may vary slightly by O/S version and vendor, but basic information should be consistent. Goes deeper into cache properties than lscpu

```
$ pwd
/sys/devices/system/cpu

$ ls
cpu0    cpu17  cpu25  cpu33  cpu41  cpu5   cpu58  cpu9    offline
cpu1    cpu18  cpu26  cpu34  cpu42  cpu50  cpu59  cpufreq online
cpu10   cpu19  cpu27  cpu35  cpu43  cpu51  cpu6   cpuidle possible
cpu11   cpu2   cpu28  cpu36  cpu44  cpu52  cpu60  hotplug power
...

$ cd cpu0/cache
$ ls
index0  index1  index2  index3  power  uevent

$ cd index0
$ ls
coherency_line_size  physical_line_partition  size
id                   power                    type
level                shared_cpu_list          uevent
number_of_sets       shared_cpu_map            ways_of_associativity
```


Expanse Cache properties – AMD Rome (AMD EPYC 7742)

level	type	line size	sets	associativity	size (KB)
L1	data	64	64	8	32
L1	instruction	64	64	8	32
L2	unified	64	1024	8	512
L3	unified	64	16384	16	16384

L1 and L2 caches are per core

L3 cache shared between 4 cores on a core complex

line size x sets x associativity = size

L2 cache size = $64 \times 1024 \times 8 = 524288 = 512 \text{ K}$

Comet Cache properties – Intel Haswell (Intel Xeon E5-2680)

level	type	line size	sets	associativity	size (KB)
L1	data	64	64	8	32
L1	instruction	64	64	8	32
L2	unified	64	512	8	256
L3	unified	64	24576	20	30720

L1 and L2 caches are per core

L3 cache shared between all 12 cores in socket

line size x sets x associativity = size

L2 cache size = $64 \times 512 \times 8 = 262144 = 256 \text{ K}$

Gordon Cache properties – Intel Sandy Bridge (Intel Xeon E5-2670)

level	type	line size	sets	associativity	size (KB)
L1	data	64	64	8	32
L1	instruction	64	64	8	32
L2	unified	64	512	8	256
L3	unified	64	16384	20	20480

L1 and L2 caches are per core

L3 cache shared between all 8 cores in socket

line size x sets x associativity = size

L2 cache size = $64 \times 512 \times 8 = 262144 = 256 \text{ K}$

Trestles Cache properties – AMD Magny-Cours (AMD Opteron Processor 6136)

level	type	line size	sets	associativity	size (KB)
L1	data	64	512	2	64
L1	instruction	64	512	2	64
L2	unified	64	512	16	512
L3	unified	64	1706	48	5118

L1 and L2 caches are per core

L3 cache shared between all 8 cores in socket

line size x sets x associativity = size

L2 cache size = $64 \times 512 \times 16 = 524288 = 512\text{K}$

Impact of cache size on performance

Note – example based on old systems, but still very illustrative

Based on the clock speed and instruction set, program run on single core of Gordon should be 2.26x faster than on Trestles. The larger L1 and L2 cache sizes on Trestles mitigate performance impact for very small problems.

DGSEV ($Ax=b$) wall times as function of problem size

N	t (Trestles)	t (Gordon)	ratio	KB
62	0.000117	0.000086	1.36	30
125	0.000531	0.000384	1.38	122
250	0.002781	0.001542	1.80	488
500	0.016313	0.007258	2.24	1953
1000	0.107222	0.046252	2.31	7812
2000	0.744837	0.331818	2.24	31250
4000	5.489990	2.464218	2.23	125000

Finding SCSI device information

SCSI (Small Computer System Interface) is a common interface for mounting peripheral, such as hard drives and SSDs. The `lsscsi` or `/proc/scsi/scsi` file command will provide info on SCSI devices.

```
$ lsscsi
[3:0:0:0]    disk      ATA          SSDSC2KB480G8R  DL67  /dev/sda
[N:0:0:1]    disk      Dell Express Flash NVMe P4510 1TB SFF__1  /dev/nvme0n1
```



Dell 1TB PCIe
NVMe Read

df provides information on filesystem usage

Local scratch (SSDs)

```
$ df -h
Filesystem                                Size  Used Avail Use% Mounted on
/dev/nvme0n1p1                            916G   77M   870G   1% /scratch
ps-071.sdsc.edu:/ps-data/community-sw      1.0T  102G   923G  10% /exppanse/community
10.21.0.21:6789,10.21.11.7:6789,10.21.11.8:6789:/ 1.7T  553G   1.2T  33% /cm/shared
10.22.101.123@o2ib:10.22.101.124@o2ib:/exppanse/projects 11P   1.4P   9.3P  13% /exppanse/lustre/projects
10.22.101.123@o2ib:10.22.101.124@o2ib:/exppanse/scratch 11P   1.4P   9.3P  13% /exppanse/lustre/scratch
10.22.100.113:/pool3/home/sinkovit         209T   4.1T  205T   2% /home/sinkovit

--- only selected filesystems shown ---
```

df provides information on filesystem usage

Community and SDSC maintained software stacks

```
$ df -h
Filesystem                                Size  Used Avail Use% Mounted on
/dev/nvme0n1p1                            916G   77M   870G   1% /scratch
ps-071.sdsc.edu:/ps-data/community-sw     1.0T  102G   923G  10% /expanse/community
10.21.0.21:6789,10.21.11.7:6789,10.21.11.8:6789:/ 1.7T  553G   1.2T  33% /cm/shared
10.22.101.123@o2ib:10.22.101.124@o2ib:/expanse/projects 11P  1.4P   9.3P  13%
/expanse/lustre/projects
10.22.101.123@o2ib:10.22.101.124@o2ib:/expanse/scratch 11P  1.4P   9.3P  13%
/expanse/lustre/scratch
10.22.100.113:/pool3/home/sinkovit        209T  4.1T  205T   2% /home/sinkovit

--- only selected filesystems shown ---
```

df provides information on filesystem usage

```
$ df -h
Filesystem                                Size  Used Avail Use% Mounted on
/dev/nvme0n1p1                            916G   77M   870G   1% /scratch
ps-071.sdsc.edu:/ps-data/community-sw      1.0T  102G   923G  10% /exppanse/community
10.21.0.21:6789,10.21.11.7:6789,10.21.11.8:6789:/ 1.7T  553G   1.2T  33% /cm/shared
10.22.101.123@o2ib:10.22.101.124@o2ib:/exppanse/projects 11P   1.4P   9.3P  13%
/exppanse/lustre/projects
10.22.101.123@o2ib:10.22.101.124@o2ib:/exppanse/scratch 11P   1.4P   9.3P  13%
/exppanse/lustre/scratch
10.22.100.113:/pool3/home/sinkovit         209T   4.1T   205T   2% /home/sinkovit

--- only selected filesystems shown ---
```

Lustre scratch and project filesystems

df provides information on filesystem usage

```
$ df -h
Filesystem                                Size  Used Avail Use% Mounted on
/dev/nvme0n1p1                            916G   77M   870G   1% /scratch
ps-071.sdsc.edu:/ps-data/community-sw      1.0T  102G   923G  10% /expanse/community
10.21.0.21:6789,10.21.11.7:6789,10.21.11.8:6789:/ 1.7T  553G   1.2T  33% /cm/shared
10.22.101.123@o2ib:10.22.101.124@o2ib:/expanse/projects 11P   1.4P   9.3P  13% /expanse/lustre/projects
10.22.101.123@o2ib:10.22.101.124@o2ib:/expanse/scratch 11P   1.4P   9.3P  13% /expanse/lustre/scratch
10.22.100.113:/pool3/home/sinkovit        209T   4.1T   205T   2% /home/sinkovit

--- only selected filesystems shown ---
```

Home filesystem

Finding network information

The ip command (/sbin/ip) is normally used by sys admins, but regular users can use it to learn about networking information

```
$ /sbin/ip link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT
group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eno1: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN mode DEFAULT group
default qlen 1000
    link/ether 6c:2b:59:bb:61:24 brd ff:ff:ff:ff:ff:ff
3: eno33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9000 qdisc mq state UP mode
DEFAULT group default qlen 1000
    link/ether 1c:34:da:62:a8:50 brd ff:ff:ff:ff:ff:ff
4: eno34: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN mode DEFAULT group
default qlen 1000
    link/ether 1c:34:da:62:a8:51 brd ff:ff:ff:ff:ff:ff
5: ib0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 4092 qdisc mq state UP mode DEFAULT
group default qlen 256
    link/infiniband 20:00:11:07:fe:80:00:00:00:00:00:00:00:00:1c:34:da:03:00:5d:53:90
brd 00:ff:ff:ff:ff:ff:12:40:1b:ff:ff:00:00:00:00:00:00:00:ff:ff:ff:ff
6: eno33.450@eno33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9000 qdisc noqueue
state UP mode DEFAULT group default qlen 1000
    link/ether 1c:34:da:62:a8:50 brd ff:ff:ff:ff:ff:ff
```


Finding OS and kernel information

Use uname to get information on the Linux kernel

```
$ uname -r
4.18.0-147.el8.x86_64
$ uname -o
GNU/Linux
$ uname -a
Linux login01 4.18.0-147.el8.x86_64 #1 SMP Wed Dec 4 21:51:45 UTC 2019 x86_64
x86_64 x86_64 GNU/Linux
```

Look in /etc/centos-release to get the Linux distribution (will vary by Linux distro)

```
$ cat /etc/centos-release
CentOS Linux release 8.1.1911 (Core)
```

Machine info - overkill?

- We've probably gone a little deeper than is necessary for you to be an effective supercomputer user.
- Think of this as a way to round out your HPC knowledge. You're learning a little bit about the tools of the trade, getting comfortable poking around on a system, acquiring the knowledge that will make it easier to work with your sys admin and picking up the background that will help you to make intelligent decisions in the future.
- Exercise: grab an interactive node (or just the login node) on Expanse and experiment with what we've covered. Cheat sheet on the next slide.
- Note that login and compute nodes are somewhat different

Machine info – cheat sheet

File or command	Information provided
<code>less /proc/cpuinfo</code> or <code>lscpu</code>	CPU specs
<code>less /proc/meminfo</code>	Memory specs and usage
<code>nvidia-smi</code>	GPU specs and usage
<code>cd /sys/devices/system/cpu/cpu0/cache</code> ... then look at directory contents	Cache configuration
<code>less /proc/scsi/scsi</code> or <code>lsscsi</code>	Peripherals (e.g. SSDs)
<code>less /etc/mtab</code>	Mounted file systems
<code>df -h</code>	File system usage (readable format)
<code>/sbin/ip link</code>	Networking information
<code>uname -a</code>	OS information
<code>less /etc/centos-release</code>	Centos version

Outline

- Expanse
 - System Overview
 - AMD EPYC Processor Architecture
- Obtaining Hardware Information
 - CPU
 - GPU
 - Memory
 - Cache
 - SCSI, df, network, OS
- top and htop tools

Using the Linux top/htop utility

The top utility is found on all Linux systems and provides a high level view of running processes. Does not give any information at the source code level (profiling), but can still be very useful for answering questions such as

- How many of my processes are running?
- What are the states of the processes (running, sleeping, etc.)?
- Which cores are being utilized?
- Are there any competing processes that may be affecting my performance?
- What fraction of the CPU is each process using?
- How much memory does each process use?
- Is the memory usage growing over time? (Useful for identifying memory leaks)
- How many threads are my processes using?

Customizing top

Top has the following defaults, but is easily customizable

- Processes only (no threads)
- To toggle threads display, type “H” while top is running
- Information for all users
- Can restrict to a single user by launching with “top -u username”
- Process ID, priority, ‘nice’ level, virtual memory, physical memory, shared memory, state, %CPU, %memory, CPU time, command
- To modify, type “f” while top is running and toggle fields using letters
- Update information every 3 seconds
- Change refresh rate by launching with “top -d *n*”
- Ordered by CPU usage
- Type “M” to order by memory usage

Non-threaded code

stivoknis — sinkovit@gcn-17-57:~ — ssh — 94x33

```
top - 08:37:00 up 60 days, 14:23, 1 user, load average: 15.32, 10.36, 6.12
Tasks: 624 total, 17 running, 607 sleeping, 0 stopped, 0 zombie
Cpu(s): 68.7%us, 1.3%sy, 0.0%ni, 29.9%id, 0.1%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 66054160k total, 37885796k used, 28168364k free, 8808k buffers
Swap: 2097144k total, 13400k used, 2083744k free, 32927192k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
70388	sinkovit	20	0	194m	76m	1612	R	100.0	0.1	1:31.06	lobfaster.pl
72547	sinkovit	20	0	120m	2976	1612	R	100.0	0.0	0:01.49	lobfaster.pl
72516	sinkovit	20	0	127m	9.9m	1608	R	100.0	0.0	0:02.09	lobfaster.pl
72526	sinkovit	20	0	121m	3388	1612	R	100.0	0.0	0:01.84	lobfaster.pl
72535	sinkovit	20	0	121m	4208	1612	R	100.0	0.0	0:01.73	lobfaster.pl
72565	sinkovit	20	0	120m	3212	1612	R	100.0	0.0	0:01.01	lobfaster.pl
72268	sinkovit	20	0	130m	12m	1612	R	98.9	0.0	0:11.96	lobfaster.pl
72359	sinkovit	20	0	123m	5976	1612	R	98.9	0.0	0:09.77	lobfaster.pl
72460	sinkovit	20	0	127m	10m	1612	R	98.9	0.0	0:08.38	lobfaster.pl
72481	sinkovit	20	0	131m	13m	1612	R	98.9	0.0	0:07.44	lobfaster.pl
72529	sinkovit	20	0	122m	4576	1612	R	98.9	0.0	0:01.82	lobfaster.pl
72439	sinkovit	20	0	130m	12m	1612	R	97.0	0.0	0:08.64	lobfaster.pl
72590	sinkovit	20	0	120m	3140	1612	R	71.7	0.0	0:00.37	lobfaster.pl
72602	sinkovit	20	0	120m	2576	1612	R	38.8	0.0	0:00.20	lobfaster.pl
72605	sinkovit	20	0	120m	2528	1600	R	34.9	0.0	0:00.18	lobfaster.pl
72608	sinkovit	20	0	119m	2340	1600	R	21.3	0.0	0:00.11	lobfaster.pl

16 processes, each using anywhere from 21.3% to 100% of a compute core.

Memory footprint (RES) is minimal, with each process only using up to 76 MB.

CPU times ranging from 0.11s (just started) to 1:31

Threaded code (thread display off)

```
stivoknis — sinkovit@gcn-17-57:~ — ssh — 87x33
Tasks: 592 total,  2 running, 590 sleeping,  0 stopped,  0 zombie
Cpu(s): 99.8%us,  0.2%sy,  0.0%ni,  0.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Mem:  66054160k total, 16519596k used, 49534564k free,   11248k buffers
Swap: 2097144k total,   13400k used, 2083744k free, 7563960k cached

  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
 81007 sinkovit  20   0 6872m 5.8g 1412 R 1595.9   9.1   5:56.48 lob_constructio
```

Threaded code with thread display toggled to the “off” position. Note the heavy CPU usage, very close to 1600%

Threaded code (thread display on)

```
stivoknis — sinkovit@gcn-17-57:~ — ssh — 87x33
Tasks: 626 total, 17 running, 609 sleeping, 0 stopped, 0 zombie
Cpu(s): 15.8%us, 0.2%sy, 0.0%ni, 84.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 66054160k total, 17495556k used, 48558604k free, 11552k buffers
Swap: 2097144k total, 13400k used, 2083744k free, 8478752k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
81007	sinkovit	20	0	6927m	5.8g	1412	R	99.5	9.2	8:37.91	lob_constructio
81096	sinkovit	20	0	6927m	5.8g	1412	R	10.5	9.2	1:13.43	lob_constructio
81105	sinkovit	20	0	6927m	5.8g	1412	R	10.5	9.2	1:13.43	lob_constructio
81107	sinkovit	20	0	6927m	5.8g	1412	R	10.5	9.2	1:13.43	lob_constructio
81097	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.40	lob_constructio
81099	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.39	lob_constructio
81100	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.44	lob_constructio
81101	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.44	lob_constructio
81102	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.43	lob_constructio
81103	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.45	lob_constructio
81106	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.44	lob_constructio
81108	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.44	lob_constructio
81109	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.29	lob_constructio
81110	sinkovit	20	0	6927m	5.8g	1412	R	10.2	9.2	1:13.39	lob_constructio
81098	sinkovit	20	0	6927m	5.8g	1412	R	9.9	9.2	1:13.44	lob_constructio
81104	sinkovit	20	0	6927m	5.8g	1412	R	9.9	9.2	1:13.38	lob_constructio

16 threads, with only one thread making good use of CPU

Total memory usage 5.8 GB (9.2% of available)

Threaded code (thread display on)

```
stivoknis — sinkovit@gcn-17-57:~ — ssh — 87x33
Tasks: 626 total, 17 running, 609 sleeping, 0 stopped, 0 zombie
Cpu(s): 90.9%us, 0.1%sy, 0.0%hi, 9.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 66054160k total, 17628152k used, 48426008k free, 11496k buffers
Swap: 2097144k total, 13400k used, 2083744k free, 8396488k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
81007	sinkovit	20	0	7132m	6.0g	1412	R	100.0	9.5	7:54.98	lob_constructio
81110	sinkovit	20	0	7132m	6.0g	1412	R	90.4	9.5	0:51.15	lob_constructio
81096	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.17	lob_constructio
81098	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.19	lob_constructio
81099	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.14	lob_constructio
81100	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.18	lob_constructio
81101	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.18	lob_constructio
81102	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.18	lob_constructio
81103	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.19	lob_constructio
81104	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.14	lob_constructio
81105	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.19	lob_constructio
81106	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.18	lob_constructio
81107	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.18	lob_constructio
81108	sinkovit	20	0	7132m	6.0g	1412	R	90.1	9.5	0:51.18	lob_constructio
81097	sinkovit	20	0	7132m	6.0g	1412	R	89.8	9.5	0:51.15	lob_constructio
81109	sinkovit	20	0	7132m	6.0g	1412	R	89.8	9.5	0:51.08	lob_constructio

16 threads, all making good (but not ideal) use of the compute cores

htop – like top,
but with better
interface, vertical
and horizontal
scrolling, process
tree view, etc.

```
1  [|||||] 100.0% 33 [|||||] 100.0% 65 [|||||] 100.0% 97 [|||||] 100.0%
2  [|||||] 100.0% 34 [|||||] 100.0% 66 [|||||] 100.0% 98 [|||||] 100.0%
3  [|||||] 100.0% 35 [|||||] 100.0% 67 [|||||] 100.0% 99 [|||||] 100.0%
4  [|||||] 100.0% 36 [|||||] 100.0% 68 [|||||] 100.0% 100 [|||||] 100.0%
5  [|||||] 100.0% 37 [|||||] 100.0% 69 [|||||] 100.0% 101 [|||||] 100.0%
6  [|||||] 100.0% 38 [|||||] 100.0% 70 [|||||] 100.0% 102 [|||||] 100.0%
7  [|||||] 100.0% 39 [|||||] 100.0% 71 [|||||] 100.0% 103 [|||||] 100.0%
8  [|||||] 100.0% 40 [|||||] 100.0% 72 [|||||] 100.0% 104 [|||||] 100.0%
9  [|||||] 100.0% 41 [|||||] 100.0% 73 [|||||] 100.0% 105 [|||||] 100.0%
10 [|||||] 100.0% 42 [|||||] 100.0% 74 [|||||] 100.0% 106 [|||||] 100.0%
11 [|||||] 100.0% 43 [|||||] 100.0% 75 [|||||] 100.0% 107 [|||||] 100.0%
12 [|||||] 100.0% 44 [|||||] 100.0% 76 [|||||] 100.0% 108 [|||||] 100.0%
13 [|||||] 100.0% 45 [|||||] 100.0% 77 [|||||] 100.0% 109 [|||||] 100.0%
14 [|||||] 100.0% 46 [|||||] 100.0% 78 [|||||] 100.0% 110 [|||||] 100.0%
15 [|||||] 100.0% 47 [|||||] 100.0% 79 [|||||] 100.0% 111 [|||||] 100.0%
16 [|||||] 100.0% 48 [|||||] 100.0% 80 [|||||] 100.0% 112 [|||||] 100.0%
17 [|||||] 100.0% 49 [|||||] 100.0% 81 [|||||] 100.0% 113 [|||||] 98.7%
18 [|||||] 100.0% 50 [|||||] 100.0% 82 [|||||] 100.0% 114 [|||||] 98.1%
19 [|||||] 100.0% 51 [|||||] 100.0% 83 [|||||] 100.0% 115 [|||||] 100.0%
20 [|||||] 100.0% 52 [|||||] 100.0% 84 [|||||] 100.0% 116 [|||||] 98.7%
21 [|||||] 99.4% 53 [|||||] 94.9% 85 [|||||] 92.9% 117 [|||||] 92.9%
22 [|||||] 100.0% 54 [|||||] 94.9% 86 [|||||] 100.0% 118 [|||||] 100.0%
23 [|||||] 99.4% 55 [|||||] 100.0% 87 [|||||] 92.9% 119 [|||||] 92.3%
24 [|||||] 100.0% 56 [|||||] 94.9% 88 [|||||] 92.3% 120 [|||||] 92.4%
25 [|||||] 99.4% 57 [|||||] 98.7% 89 [|||||] 98.1% 121 [|||||] 98.7%
26 [|||||] 100.0% 58 [|||||] 98.7% 90 [|||||] 100.0% 122 [|||||] 98.7%
27 [|||||] 99.4% 59 [|||||] 100.0% 91 [|||||] 98.7% 123 [|||||] 98.7%
28 [|||||] 98.7% 60 [|||||] 98.7% 92 [|||||] 98.7% 124 [|||||] 100.0%
29 [|||||] 100.0% 61 [|||||] 100.0% 93 [|||||] 100.0% 125 [|||||] 100.0%
30 [|||||] 100.0% 62 [|||||] 100.0% 94 [|||||] 100.0% 126 [|||||] 100.0%
31 [|||||] 100.0% 63 [|||||] 100.0% 95 [|||||] 100.0% 127 [|||||] 100.0%
32 [|||||] 100.0% 64 [|||||] 100.0% 96 [|||||] 100.0% 128 [|||||] 100.0%
Mem[|||||] 231G/251G
Swp[|||||] 0K/0K
Tasks: 93, 441 thr; 128 running
Load average: 107.02 47.30 22.84
Uptime: 89 days, 21:28:38
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
85153	sinkovit	20	0	8061M	7024M	21896	R	398.	2.7	7:58.93	xhpl
85145	sinkovit	20	0	8061M	7024M	21836	R	394.	2.7	7:57.31	xhpl
85149	sinkovit	20	0	8003M	6966M	22000	R	398.	2.7	7:58.78	xhpl
85157	sinkovit	20	0	8003M	6967M	22040	R	398.	2.7	8:01.77	xhpl
85137	sinkovit	20	0	8061M	7024M	21904	R	398.	2.7	7:51.81	xhpl

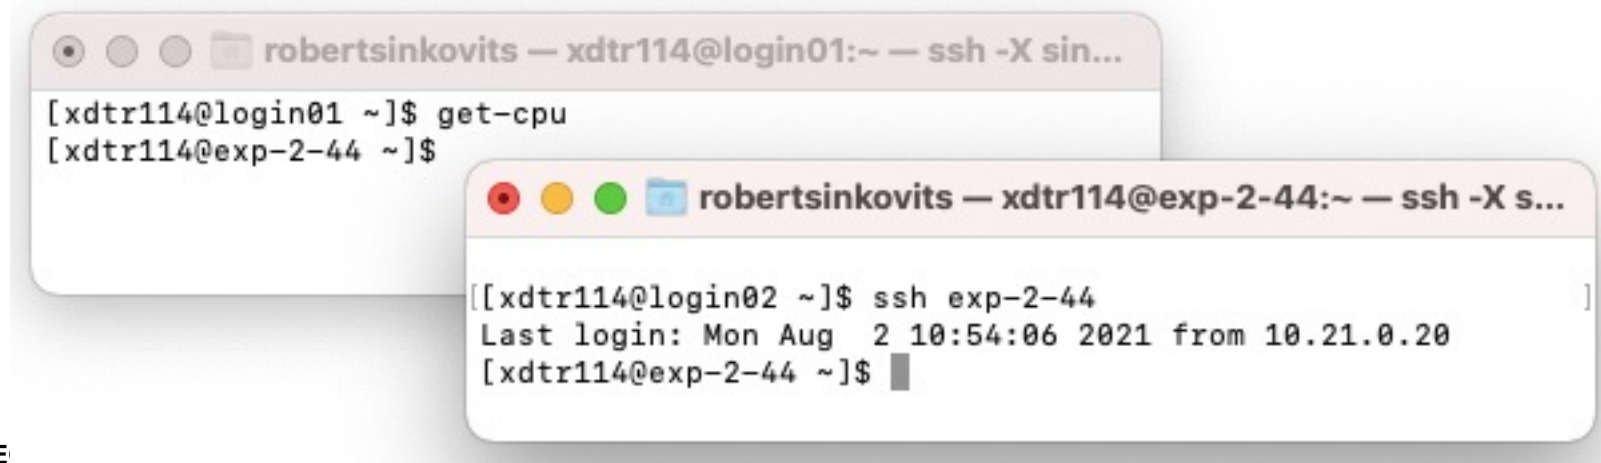
htop – like top,
but with better
interface, vertical
and horizontal
scrolling, process
tree view, etc.

1	[]	89.8%	33	[]	82.1%	65	[]	86.5%	97	[]	87.7%
2	[]	87.2%	34	[]	94.2%	66	[]	81.9%	98	[]	94.2%
3	[]	74.4%	35	[]	74.8%	67	[]	82.5%	99	[]	79.5%
4	[]	74.7%	36	[]	75.6%	68	[]	94.9%	100	[]	78.7%
5	[]	92.3%	37	[]	91.7%	69	[]	93.6%	101	[]	93.6%
6	[]	85.3%	38	[]	99.4%	70	[]	94.8%	102	[]	94.8%
7	[]	94.2%	39	[]	87.9%	71	[]	86.5%	103	[]	85.8%
8	[]	84.5%	40	[]	90.4%	72	[]	87.8%	104	[]	85.4%
9	[]	91.7%	41	[]	94.2%	73	[]	92.9%	105	[]	94.8%
10	[]	84.0%	42	[]	86.5%	74	[]	85.3%	106	[]	85.3%
11	[]	83.3%	43	[]	91.0%	75	[]	85.3%	107	[]	85.3%
12	[]	92.9%	44	[]	86.4%	76	[]	92.9%	108	[]	91.6%
13	[]	87.1%	45	[]	86.5%	77	[]	100.0%	109	[]	96.8%
14	[]	99.4%	46	[]	97.4%	78	[]	90.4%	110	[]	94.3%
15	[]	87.8%	47	[]	90.3%	79	[]	88.5%	111	[]	91.6%
16	[]	85.9%	48	[]	87.7%	80	[]	89.0%	112	[]	91.6%
17	[]	60.0%	49	[]	55.8%	81	[]	61.1%	113	[]	62.6%
18	[]	56.4%	50	[]	99.4%	82	[]	58.4%	114	[]	59.0%
19	[]	54.2%	51	[]	54.8%	83	[]	98.7%	115	[]	98.7%
20	[]	98.7%	52	[]	55.8%	84	[]	58.3%	116	[]	59.6%
21	[]	68.4%	53	[]	67.3%	85	[]	62.2%	117	[]	65.0%
22	[]	89.1%	54	[]	55.8%	86	[]	61.9%	118	[]	59.6%
23	[]	54.8%	55	[]	92.3%	87	[]	59.4%	119	[]	96.2%
24	[]	54.5%	56	[]	56.7%	88	[]	98.7%	120	[]	60.6%
25	[]	74.2%	57	[]	92.2%	89	[]	82.6%	121	[]	82.6%
26	[]	62.4%	58	[]	58.4%	90	[]	73.5%	122	[]	94.9%
27	[]	90.4%	59	[]	57.1%	91	[]	92.9%	123	[]	74.8%
28	[]	61.5%	60	[]	68.4%	92	[]	73.7%	124	[]	74.7%
29	[]	87.7%	61	[]	88.5%	93	[]	96.8%	125	[]	90.4%
30	[]	99.4%	62	[]	86.6%	94	[]	96.2%	126	[]	98.7%
31	[]	88.4%	63	[]	98.1%	95	[]	89.8%	127	[]	93.5%
32	[]	88.4%	64	[]	87.2%	96	[]	91.1%	128	[]	91.7%
Mem	[]				231G/251G	Tasks: 93, 433 thr; 128 running					
Swp	[]				0K/0K	Load average: 93.28 37.71 18.93					
Uptime: 89 days, 21:28:04											

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
85156	sinkovit	20	0	7937M	6900M	21844	R	371.	2.7	5:50.21	xhpl
85160	sinkovit	20	0	8007M	6970M	21880	R	369.	2.7	5:49.53	xhpl
85152	sinkovit	20	0	8007M	6970M	22048	R	367.	2.7	5:47.94	xhpl
85138	sinkovit	20	0	8006M	6969M	21948	R	364.	2.7	5:43.53	xhpl
85148	sinkovit	20	0	7937M	6901M	22112	R	362.	2.7	5:46.87	xhpl

top/htop exercise

- Once we do the hands on exercises for the other sessions, open another terminal and run top or htop
- In the meantime, try running on the login node or an interactive node and play around with the options
- Useful tip – on Expanse, once Slurm has allocated a node to your account, you can then directly login to the node



The image shows two overlapping terminal windows. The top window has a title bar 'robertsinkovits — xdtr114@login01:~ — ssh -X sin...' and contains the following text:
[xdtr114@login01 ~]\$ get-cpu
[xdtr114@exp-2-44 ~]\$

The bottom window has a title bar 'robertsinkovits — xdtr114@exp-2-44:~ — ssh -X s...' and contains the following text:
[[xdtr114@login02 ~]\$ ssh exp-2-44
Last login: Mon Aug 2 10:54:06 2021 from 10.21.0.20
[xdtr114@exp-2-44 ~]\$